



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

AdaBoost

Elie Kawerk
Data Scientist



Boosting

- **Boosting:** Ensemble method combining several weak learners to form a strong learner.
- **Weak learner:** Model doing slightly better than random guessing.
- Example of weak learner: Decision stump (CART whose maximum depth is 1).



Boosting

- Train an ensemble of predictors sequentially.
- Each predictor tries to correct its predecessor.
- Most popular boosting methods:
 - AdaBoost,
 - Gradient Boosting.




Adaboost

- Stands for **Adaptive Boosting**.
- Each predictor pays more attention to the instances wrongly predicted by its predecessor.
- Achieved by changing the weights of training instances.
- Each predictor is assigned a coefficient α .
- α depends on the predictor's training error.




AdaBoost: Training





Learning Rate

Learning rate: $0 < \eta \leq 1$



{{1}}



AdaBoost: Prediction

- Classification:
 - Weighted majority voting.
 - In `sklearn`: `AdaBoostClassifier`.
- Regression:
 - Weighted average.
 - In `sklearn`: `AdaBoostRegressor`.

AdaBoost Classification in sklearn (Breast Cancer dataset)

```
# Import models and utility functions
In [1]: from sklearn.ensemble import AdaBoostClassifier
In [2]: from sklearn.tree import DecisionTreeClassifier
In [3]: from sklearn.metrics import roc_auc_score
In [4]: from sklearn.model_selection import train_test_split

# Set seed for reproducibility
In [5]: SEED = 1

# Split data into 70% train and 30% test
In [6]: X_train, X_test, y_train, y_test = \
        train_test_split(X, y,
                        test_size=0.3,
                        stratify=y,
                        random_state=SEED)
```


AdaBoost Classification in sklearn (Breast Cancer dataset)

```
# Instantiate a classification-tree 'dt'
In [7]: dt = DecisionTreeClassifier(max_depth=1,
                                   random_state=SEED)

# Instantiate an AdaBoost classifier 'adb_clf'
In [8]: adb_clf = AdaBoostClassifier(base_estimator=dt,
                                     n_estimators=100)

# Fit 'adb_clf' to the training set
In [9]: adb_clf.fit(X_train, y_train)

# Predict the test set probabilities of positive class
In [10]: y_pred_proba = adb_clf.predict_proba(X_test)[:,1]

# Evaluate test-set roc_auc_score
In [11]: adb_clf_roc_auc_score = roc_auc_score(y_test, y_pred_proba)
```



AdaBoost Classification in sklearn (Breast Cancer dataset)

```
# Print adb_clf_roc_auc_score
In [12]: print('ROC AUC score: {:.2f}'.format(adb_clf_roc_auc_score))

Out[12]: ROC AUC score: 0.99
```



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Let's practice!



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Gradient Boosting (GB)

Elie Kawerk

Data Scientist




Gradient Boosted Trees

- Sequential correction of predecessor's errors.
- Does not tweak the weights of training instances.
- Fit each predictor is trained using its predecessor's residual errors as labels.
- Gradient Boosted Trees: a CART is used as a base learner.




Gradient Boosted Trees for Regression: Training





Shrinkage





Gradient Boosted Trees: Prediction

- Regression:
 - $y_{pred} = y_1 + \eta r_1 + \dots + \eta r_N$
 - In `sklearn`: `GradientBoostingRegressor`.
- Classification:
 - In `sklearn`: `GradientBoostingClassifier`.

Gradient Boosting in sklearn (auto dataset)

```
# Import models and utility functions
In [1]: from sklearn.ensemble import GradientBoostingRegressor
In [2]: from sklearn.model_selection import train_test_split
In [3]: from sklearn.metrics import mean_squared_error as MSE

# Set seed for reproducibility
In [4]: SEED = 1

# Split dataset into 70% train and 30% test
In [5]: X_train, X_test, y_train, y_test = \
        train_test_split(X, y,
                        test_size=0.3,
                        random_state=SEED)
```



Gradient Boosting in sklearn (auto dataset)

```
# Instantiate a GradientBoostingRegressor 'gbt'
In [6]: gbt = GradientBoostingRegressor(n_estimators=300,
                                         max_depth=1,
                                         random_state=SEED)

# Fit 'gbt' to the training set
In [7]: gbt.fit(X_train, y_train)

# Predict the test set labels
In [8]: y_pred = gbt.predict(X_test)

# Evaluate the test set RMSE
In [9]: rmse_test = MSE(y_test, y_pred)**(1/2)

# Print the test set RMSE
In [10]: print('Test set RMSE: {:.2f}'.format(rmse_test))

Out[10]: Test set RMSE: 4.01
```



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Let's practice!



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Stochastic Gradient Boosting (SGB)

Elie Kawerk

Data Scientist



Gradient Boosting: Cons

- GB involves an exhaustive search procedure.
- Each CART is trained to find the best split points and features.
- May lead to CARTs using the same split points and maybe the same features.




Stochastic Gradient Boosting

- Each tree is trained on a random subset of rows of the training data.
- The sampled instances (40%-80% of the training set) are sampled without replacement.
- Features are sampled (without replacement) when choosing split points.
- Result: further ensemble diversity.
- Effect: adding further variance to the ensemble of trees.



Stochastic Gradient Boosting: Training





Stochastic Gradient Boosting in sklearn (auto dataset)

```
# Import models and utility functions
In [1]: from sklearn.ensemble import GradientBoostingRegressor
In [2]: from sklearn.model_selection import train_test_split
In [3]: from sklearn.metrics import mean_squared_error as MSE

# Set seed for reproducibility
In [4]: SEED = 1

# Split dataset into 70% train and 30% test
In [5]: X_train, X_test, y_train, y_test = \
        train_test_split(X, y,
                        test_size=0.3,
                        random_state=SEED)
```


Stochastic Gradient Boosting in sklearn (auto dataset)

```
# Instantiate a stochastic GradientBoostingRegressor 'sgbt'
In [6]: sgbt = GradientBoostingRegressor(max_depth=1,
                                         subsample=0.8,
                                         max_features=0.2,
                                         n_estimators=300,
                                         random_state=SEED)

# Fit 'sgbt' to the training set
In [7]: sgbt.fit(X_train, y_train)

# Predict the test set labels
In [8]: y_pred = sgbt.predict(X_test)
```



Stochastic Gradient Boosting in sklearn (auto dataset)

```
# Evaluate test set RMSE 'rmse_test'
In [9]: rmse_test = MSE(y_test, y_pred)**(1/2)

# Print 'rmse_test'
In [10]: print('Test set RMSE: {:.2f}'.format(rmse_test))

Out[10]: Test set RMSE: 3.95
```



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Let's practice!